

Runtime File Edit View

Run all Kernel

Compute not connected



Exercise: More metrics derived from confusion matrices

In this exercise, we'll learn about different metrics, using them to explain the results obtained from the *binary classification model* we built in the previous unit.

Data visualization

We'll use the dataset with different classes of objects found on the mountain one more time:

```
import pandas
import numpy
!wget https://raw.githubusercontent.com/MicrosoftDocs/mslearn-introduction-to-machine-learning
!wget https://raw.githubusercontent.com/MicrosoftDocs/mslearn-introduction-to-machine-learning

#Import the data from the .csv file
dataset = pandas.read_csv('snow_objects.csv', delimiter="\t")

#Let's have a look at the data
dataset
```

Recall that to use the dataset above for *binary classification*, we need to add another column to the dataset, and set it to `True` where the original label is `hiker` and `False` where it's not

where the original label is `hiker`, and `False` where it's not.

Let's then add that label, split the dataset and train the model again:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split

# Add a new label with true/false values to our dataset
dataset["is_hiker"] = dataset.label == "hiker"

# Split the dataset in an 70/30 train/test ratio.
train, test = train_test_split(dataset, test_size=0.3, random_state=1, shuffle=True)

# define a random forest model
model = RandomForestClassifier(n_estimators=1, random_state=1, verbose=False)

# Define which features are to be used
features = ["size", "roughness", "motion"]

# Train the model using the binary label
model.fit(train[features], train.is_hiker)

print("Model trained!")
```

We can now use this model to predict whether objects in the snow are hikers or not.

Let's plot its *confusion matrix*:

```
# sklearn has a very convenient utility to build confusion matrices
from sklearn.metrics import confusion_matrix
import plotly.figure_factory as ff

# Calculate the model's accuracy on the TEST set
actual = test.is_hiker
```

```
predictions = model.predict(test[features])

# Build and print our confusion matrix, using the actual values and predictions
# from the test set, calculated in previous cells
cm = confusion_matrix(actual, predictions, normalize=None)

# Create the list of unique labels in the test set, to use in our plot
# I.e., ['True', 'False',]
unique_targets = sorted(list(test["is_hiker"].unique()))

# Convert values to lower case so the plot code can count the outcomes
x = y = [str(s).lower() for s in unique_targets]

# Plot the matrix above as a heatmap with annotations (values) in its cells
fig = ff.create_annotated_heatmap(cm, x, y)

# Set titles and ordering
fig.update_layout( title_text="<b>Confusion matrix</b>",
                   yaxis = dict(categoryorder = "category descending"))

fig.add_annotation(dict(font=dict(color="black",size=14),
                        x=0.5,
                        y=-0.15,
                        showarrow=False,
                        text="Predicted label",
                        xref="paper",
                        yref="paper"))

fig.add_annotation(dict(font=dict(color="black",size=14),
                        x=-0.15,
                        y=0.5,
                        showarrow=False,
                        text="Actual label",
                        textangle=-90,
                        xref="paper",
                        yref="paper"))

# We need margins so the titles fit
fig.update_layout(margin=dict(t=80, r=20, l=120, b=50))
fig['data'][0]['showscale'] = True
fig.show()
```

```
# Let's also calculate some values that will be used throughout this exercise
# We already have actual values and corresponding predictions, defined above
correct = actual == predictions
tp = numpy.sum(correct & actual)
tn = numpy.sum(correct & numpy.logical_not(actual))
fp = numpy.sum(numpy.logical_not(correct) & actual)
fn = numpy.sum(numpy.logical_not(correct) & numpy.logical_not(actual))

print("TP - True Positives: ", tp)
print("TN - True Negatives: ", tn)
print("FP - False positives: ", fp)
print("FN - False negatives: ", fn)
```

We can use the preceding values and matrix to help us understand other metrics.

Calculating metrics

From here on, we'll take a closer look at each at the following metrics, how they're calculated, and how they can help explain our current model.

- Accuracy
- Sensitivity/Recall
- Specificity
- Precision
- False positive rate

Let's first recall some useful terms:

- TP = True positives: a positive label is correctly predicted
- TN = True negatives: a negative label is correctly predicted
- FP = False positives: a negative label is predicted as a positive

- FN = False negatives: a positive label is predicted as a negative

Accuracy

Accuracy is the number of correct predictions divided by the total number of predictions:

$$\text{accuracy} = (\text{TP} + \text{TN}) / \text{number of samples}$$

It's possibly the most basic metric used but, as we've seen, it's not the most reliable when *imbalanced datasets* are used.

In code:

```
# Calculate accuracy
# len(actual) is the number of samples in the set that generated TP and TN
accuracy = (tp+tn) / len(actual)

# print result as a percentage
print(f"Model accuracy is {accuracy:.2f}%")
```

Sensitivity/Recall

Sensitivity and *Recall* are interchangeable names for the same metric, which expresses the fraction of samples **correctly** predicted by a model:

$$\text{sensitivity} = \text{recall} = \text{TP} / (\text{TP} + \text{FN})$$

This is an important metric, that tells us how out of all the *actually positive* samples, how many are **correctly** predicted as positive.

In code:

```
# code for sensitivity/recall
```

```
sensitivity = recall = tp / (tp + fn)

# print result as a percentage
print(f"Model sensitivity/recall is {sensitivity:.2f}%")
```

Specificity

Specificity expresses the fraction of **negative** labels correctly predicted over the total number of existing negative samples:

$$\text{specificity} = \text{TN} / (\text{TN} + \text{FP})$$

Specificity tells us how out of all the *actually* **negative** samples, how many are **correctly** predicted as negative.

We can calculate it using the following code:

```
# Code for specificity
specificity = tn / (tn + fp)

# print result as a percentage
print(f"Model specificity is {specificity:.2f}%")
```

Precision

Precision expresses the proportion of **correctly** predicted positive samples over all positive predictions:

$$\text{precision} = \text{TP} / (\text{TP} + \text{FP})$$

In other words, it indicates how out of all positive predictions, how many are truly positive labels.

We can calculate it using the following code:

```
# Code for precision

precision = tp / (tp + fp)

# print result as a percentage
print(f"Model precision is {precision:.2f}%")
```

False positive rate

False positive rate or FPR, is the number of **incorrect** positive predictions divided by the total number of negative samples:

$$\text{false_positive_rate} = \text{FP} / (\text{FP} + \text{TN})$$

```
# Code for false positive rate
false_positive_rate = fp / (fp + tn)

# print result as a percentage
print(f"Model false positive rate is {false_positive_rate:.2f}%")
```

Notice that the sum of `specificity` and `false positive rate` should always be equal to 1 .

Conclusion

There are several different metrics that can help us evaluate the performance of a model in the context of the quality of its predictions.

The choice of the most adequate metrics, however, is primarily a function of the data and the problem we're trying to solve.

Summary

We covered the following topics in this unit:

 No compute  Compute not connected  Viewing

Kernel not connected

Next unit: Knowledge check

Continue >

How are we doing? ☆ ☆ ☆ ☆ ☆